

## Overview

A **B-tree** is a self-balancing tree data structure that maintains sorted data and allows searches, sequential access, insertions, and deletions in logarithmic time. It is commonly used in databases and file systems. The B-tree is characterized by its ability to store multiple keys in a single node, with each node potentially having many children. This structure helps reduce the height of the tree, improving efficiency for large datasets. B-trees are especially useful for systems that read and write large blocks of data.

## Why B-tree

B-trees are widely used because they efficiently manage large amounts of data by maintaining balance, which ensures that operations like search, insert, delete, and sequential access can be performed in logarithmic time. They are particularly useful in databases and file systems where data is stored on disk, as B-trees minimize the number of disk reads by allowing more data to be stored in each node. This makes B-trees ideal for handling large datasets with frequent read and write operations, ensuring performance remains consistent.

## How B-tree achieves the speed

A B-tree is fast because it minimizes the height of the tree, ensuring that data is accessed with a minimal number of disk reads. Each node in a B-tree can store multiple keys (multilevel indexing) and have multiple children, which allows more data to be kept in memory and reduces the number of I/O operations needed to locate or modify data.

## What is multilevel indexing

**Multilevel indexing** is a technique used to manage large datasets in databases, where single-level indexing may not be efficient. It involves creating a hierarchy of indexes to reduce the number of disk accesses needed to find data.

Here's how it works:

1. **Primary Index:** The first level of indexing, often pointing to blocks of data.
2. **Secondary Index:** An index on the primary index, further reducing the search space.
3. **Third (or more) Levels:** Additional indexes as needed, forming a tree-like structure.

This hierarchical approach improves the speed of data retrieval by reducing the number of I/O operations, making it more efficient for large-scale data systems.

### Compare BST, Unbalanced BST, Sorted Array and B-tree for operations

Here's a comparison of the number of comparisons needed for insert, delete, and search operations in different data structures for 1 million records :

	Search	Insert	Delete
Balanced BST	20	20	20
Unbalanced BST	$10^6$	$10^6$	$10^6$
Sorted Array	20	$10^6$	$10^6$
B-Tree (Order 100)	3	3	3

#### 1. Sorted Array:

Search:  $O(\log n)$  using binary search.

Insert:  $O(n)$  because elements must be shifted.

Delete:  $O(n)$  for the same reason as insertion.

#### 2. Unbalanced Binary Search Tree (BST):

Search:  $O(n)$  in the worst case (e.g., skewed tree).

Insert:  $O(n)$  in the worst case.

Delete:  $O(n)$  in the worst case.

#### 3. Balanced Binary Search Tree (e.g., AVL, Red-Black Tree):

Search:  $O(\log n)$

Insert:  $O(\log n)$

Delete:  $O(\log n)$

#### 4. B-tree:

Search:  $O(\log n)$ .

Insert:  $O(\log n)$ .

Delete:  $O(\log n)$ .

## Summary:

Sorted Arrays offer efficient search but slow insert/delete operations due to shifting. Unbalanced BSTs can degrade to linear time if not balanced. Balanced BSTs and B-trees maintain logarithmic time complexity across all operations, with B-trees particularly optimized for disk-based storage systems.

## Why B-tree is also known as M-way tree

A B-tree is also known as an **M-way tree** because it allows each node to have multiple children (up to M), where M is the order of the tree. Unlike binary trees, which are limited to two children per node.

## Structure of B-tree

The structure of a B-tree is designed to maintain balance and efficiency for large datasets. Here's a breakdown:

1. **Nodes:** Each node can contain multiple keys (up to  $M-1$  keys for a B-tree of order  $M$ ) and multiple child pointers (up to  $M$  children).
2. **Root Node:** The top node, which can have between 1 and  $M-1$  keys.
3. **Internal Nodes:** Nodes between the root and leaves, containing keys that partition the data and pointers to child nodes.
4. **Leaf Nodes:** The bottom-most nodes that contain keys and may point to the actual data records.
5. **Balanced Structure:** The tree remains balanced, ensuring all leaf nodes are at the same level.
6. **Key Distribution:** Keys within a node are sorted, and each child pointer separates ranges of keys.

This structure allows for efficient searching, insertion, and deletion operations, with the tree growing in a balanced manner.

## Characteristics of B-tree

The main characteristics of a B-tree are:

1. **Balanced Tree:** All leaf nodes are at the same level, ensuring balanced height.

2. **Multiple Keys and Children:** Each node can hold multiple keys and have multiple child pointers.
3. **Order:** Defined by order  $M$ , where each node has at most  $M-1$  keys and  $M$  children.
4. **Efficient Operations:** Supports efficient search, insertion, and deletion operations in  $O(\log n)$  time.
5. **Disk Optimization:** Minimizes disk I/O operations by keeping nodes large enough to fit into a single disk page.

## Find height of b-tree

The height “ $h$ ” of a B-tree of order “ $M$ ” with “ $n$ ” keys can be estimated using the following formula:

$$h \leq \log_{\lceil M/2 \rceil} \left( \frac{n+1}{2} \right)$$

Explanation:

- “ $M$ ” is the order of the B-tree (maximum number of children per node).
- “ $n$ ” is the total number of keys in the tree.
- The height “ $h$ ” is the maximum number of levels in the tree.

## B-tree vs B+ tree

Here’s a comparison between **B-tree** and **B+ tree**:

### B-tree:

**Data Storage:** Keys and data are stored in both internal and leaf nodes.

**Search Efficiency:** Searching can be slightly slower because data is spread across all nodes.

**Structure:** More balanced but less optimized for range queries.

**In-order Traversal:** Difficult, as not all nodes are linked.

B+ tree:

**Data Storage:** All data is stored only in leaf nodes; internal nodes only store keys.

**Search Efficiency:** Faster for range queries because leaves are linked.

**Structure:** Optimized for sequential access and range queries.

**In-order Traversal:** Easier, as all leaf nodes are linked in a sequential manner.

**Summary:** B+ trees are more efficient for range queries and sequential data access, while B-trees are slightly more balanced for general-purpose searches.